



Course name: Mathematical Software Programming

Aids allowed: All aids allowed

Exam duration: 4 hours

Weighting: 80/100

Final exam Mathematical Software Programming

This exam contains a total of 18 questions: 14 multiple choice questions (questions 1–14) and 4 programming questions (questions 15–18). Your exam answers must be submitted electronically as a **PDF document**. You may include your code in the document along with your answers or submit the code separately in a **ZIP file**.

1. (4 points) Properties of floating-point arithmetic.

- (a) The associative property of multiplication holds for floating-point arithmetic, *i.e.*,

$$a(bc) = (ab)c$$

where a , b , and c are floating-point numbers.

A. True

B. False

- (b) The distributive property of multiplication holds for floating-point arithmetic, *i.e.*,

$$a(b + c) = ab + ac$$

where a , b , and c are floating-point numbers.

A. True

B. False

2. (2 points) What is the condition number of $f(x) = \cos(x)$?

A. $|x \tan(x)|$

B. $|\tan(x)|$

C. $|\sin(x)|$

D. $|x \sin(x)|$

$$\left| \frac{x f'(x)}{f(x)} \right|$$

3. (2 points) Given n numbers x_1, x_2, \dots, x_n where $\sum_{i=1}^n |x_i| \approx 500$ and $|\sum_{i=1}^n x_i| \approx 0.01$, what is approximately the worst-case relative error for sequential floating-point summation of the numbers if the unit round-off is u and $nu < 0.1$?

A. $5.3nu$

B. $53nu$

C. $530nu$

D. $5300nu$

E. $53000nu$

$$\frac{|\hat{S}_n - S_n|}{|S_n|} \leq c \cdot 1.06 nu$$

$$c = \frac{\sum_{i=1}^n |x_i|}{|\sum_{i=1}^n x_i|}$$

4. (4 points) A programmer wrote the following C function for evaluating quadratic functions of the form

$$f(x) = x^T A x$$

where x is a vector of length n and A is a symmetric matrix of order n :

```
double quad_form(double **A, double *x, unsigned int n) {  
    unsigned int i,j;  
    double res = 0.0;  
    double *y = malloc(n*sizeof(*y));  
    if (y==NULL) return NAN;  
    for (i=0;i<n;i++) {  
        y[i] = 0.0;  
        for (j=0;j<n;j++)  
            y[i] += A[i][j]*x[j];  
        res += y[i]*x[i];  
    }  
    return res;  
}
```

There is a serious problem with this implementation. What is it?

- ☒ A. The function returns `res` which is a local variable.
 - ☐ B. The function does not give the correct answer.
 - ☒ C. The function leaks memory.
 - ☐ D. The array `y` is too small, leading to undefined behavior.
 - ☐ E. The array `y` is not properly initialized.
5. (2 points) Improved locality generally leads to
- ☐ A. improved numerical accuracy
 - ☐ B. worse numerical accuracy
 - ☐ C. fewer cache hits
 - ☒ D. fewer cache misses

6. (4 points) The theoretical improvement in speed of execution of a task executed on p processors can be expressed as

$$S(p) = \frac{T(1)}{T(p)} = \frac{\cancel{fT(1)} + (1 - \cancel{f})T(1)}{(f/p)\cancel{T(1)} + (1 - \cancel{f})\cancel{T(1)}}$$

where $T(p)$ is the execution time on p processors (real time) and f is the so-called parallel fraction of the task. For example, if 50% of a task can be parallelized, then $f = 0.5$. Suppose we measure $T(1) = 100$ and $T(4) = 37$.

- (a) What is the value of the parallel fraction f of the task?

A. $f = 0.37$

B. $f = 0.71$

C. $f = 0.84$

D. $f = 0.91$

$$\frac{37}{100} = \frac{f}{p} + 1 - f$$

- (b) What is the theoretical execution time on 8 processors?

A. $T(8) = 20.4$

B. $T(8) = 26.5$

C. $T(8) = 29.0$

D. $T(8) = 37.9$

7. (4 points) The real-valued function

$$f(x) = \frac{ax + b}{cx + d}$$

is a so-called linear fractional function. We will assume that $ad - bc \neq 0$ and $c \neq 0$.

- (a) What is the condition number of f ?

A. $\left| \frac{x(ad-bc)}{ac(x+b/a)(x+d/c)} \right|$

B. $\left| \frac{(ad-bc)}{ac(x+b/a)(x+d/c)} \right|$

C. $\left| \frac{x}{(x+b/a)(x+d/c)} \right|$

D. $\left| \frac{(ad-bc)}{ac(x+d/c)} \right|$

$$\left| \frac{x \cdot f'(x)}{f(x)} \right|$$

- (b) Catastrophic cancellation is likely to occur when

A. $ax + b \approx cx + d$

B. $ax + b \approx bx + a$

C. $x \approx a/b$ or $x \approx c/d$

D. $x \approx -b/a$ or $x \approx -d/c$

8. (4 points) A half-precision floating point number occupies 16 bits and has the following representation

s	$e_1 \dots e_5$	$d_1 d_2 \dots d_{10}$
-----	-----------------	------------------------

where s is the sign bit, d_i is the i th bit of the mantissa, and e_i is the i th bit of the exponent. Thus, a half-precision floating point number can be represented as

$$x = (-1)^s \cdot (d_0.d_1d_2\dots,d_{10})_2 \cdot 2^E = (-1)^s \cdot \sum_{i=0}^{10} d_i 2^{E-i}$$

where $E \in \{-14, -13, \dots, 14, 15\}$ is a decimal representation of the exponent.

- (a) What is the unit round-off for the half-precision floating point format?

A. $u = 2^{-10}$

B. $u = 2^{-11}$

C. $u = 2^{-12}$

D. $u = 2^{-13}$

- (b) Recall the following relative error bound for the sequential summation of n floating-point numbers x_1, \dots, x_n

$$\epsilon_{\text{rel}} \leq \frac{\sum_{i=1}^n |x_i|}{|\sum_{i=1}^n x_i|} \cdot 1.06 \cdot nu, \quad nu < 0.1.$$

Now suppose x_1, \dots, x_n are nonnegative numbers. What is approximately the largest value of n for which we can guarantee a relative error of at most 10% using half-precision arithmetic?

A. $n \approx 96$

B. $n \approx 102$

C. $n \approx 193$

D. $n \approx 205$

$$1 \cdot 1.06 \cdot nu \leq 0.1$$

$$n \leq \frac{0.1}{1.06} \cdot 2^{11}$$


2048

9. (6 points) Given a vector $x = (x_1, \dots, x_n)$ of length n , we can compute the sum $s_n = \sum_{i=1}^n x_i$ recursively using the following C function:

```
double rsum(double *x, unsigned int n) {  
    unsigned int m = n/2;  
    if (m>1)  
        return rsum(x,m)+rsum(x+m,n-m);  
    else if (m==1 && n-m>1)  
        return x[0]+rsum(x+m,n-m);  
    else if (m==1 && n-m==1)  
        return x[0]+x[1];  
    else  
        return x[0];  
}
```

- (a) What type of recursion is this?
- A. Single recursion
 - B. Multiple recursion
- (b) What is the time complexity of computing s_n using `rsum(x,n)`?
- A. $O(1)$
 - B. $O(\log(n))$
 - C. $O(n)$
 - D. $O(2^n)$
- (c) Excluding the space required to store x , what is the space complexity of computing s_n using `rsum(x,n)`?
- A. $O(1)$
 - B. $O(\log(n))$
 - C. $O(n)$
 - D. $O(2^n)$

10. (4 points) The following plot shows the CPU time required by some algorithm to solve a certain problem as a function of its dimension n .



What is the time complexity of the algorithm?

- A. $O(1)$
 - B. $O(n)$
 - C. $O(n^2)$
 - D. $O(n^3)$
11. (2 points) The purpose of loop unrolling is to
- A. reduce execution time
 - B. improve numerical accuracy
 - C. improve code readability
 - D. avoid while loops
12. (2 points) In the C++ programming language, the datatype `int&` is
- A. the address of an `int`
 - B. a pointer to an `int`
 - C. a reference to an `int`
 - D. the size of an `int` in bytes

13. (2 points) The `std::vector` class template is a container that represents a
- A. linked list
 - B. stack
 - C. static array
 - D. dynamic array
14. (4 points) A programmer wrote the following function to check if two vectors x and y of length n are orthogonal within numerical precision:

```
int orthogonal(unsigned int n, double *x, double *y) {  
    double dot = 0.0;  
    for (unsigned int i=0; i<n; i++)  
        dot += x[i]*y[i];  
    if (fabs(dot) <= 1.06*n*DBL_EPSILON)  
        return 1;    // x and y are orthogonal  
    else  
        return 0;    // x and y are not orthogonal  
}
```

Assuming that all floating-point numbers are normal and that there is no under- or overflow, will this implementation work as intended?

- A. Yes
 - B. No
15. (8 points) Write a C function that evaluates the function

$$f(n) = \log(n!)$$

with domain \mathbb{N}_0 (the set of nonnegative integers).

Hint: We have that $0! = 1$.

Use the following function prototype:

```
double logfactorial(unsigned int n);
```

Test your code with $n \in \{0, 1, 2, 5, 50, 100, 1000\}$.

16. (8 points) Given n linear functions of $x \in \mathbb{R}$

$$g_i(x) = a_i x + b_i, \quad i = 1, \dots, n,$$

we seek to compute the maximum of the n functions, *i.e.*,

$$f(x) = \max_{i=1, \dots, n} (a_i x + b_i).$$

Implement a function that evaluates $f(x)$ given x and two length n arrays with the coefficients a_1, \dots, a_n and b_1, \dots, b_n .

Use the following prototype:

```
double linear_max(  
    const double x,  
    const double *a, /* array with a coefficients */  
    const double *b, /* array with b coefficients */  
    unsigned int n  
);
```

17. (8 points) Given a nonzero vector $v \in \mathbb{R}^n$, the matrix

$$T = I - 2 \frac{vv^T}{\|v\|_2^2}$$

is a so-called Householder matrix.

- (a) Write a function that takes a vector x and a vector v as inputs (both of length n) and computes $x := Tx$ (*i.e.*, the function should overwrite x with the result Tx). The function should have the following prototype:

```
int householder(const double *v, double *x, unsigned int n);
```

The return value should be zero if successful. In case of an error, the function should return a nonzero value and the array x should remain unchanged.

- (b) What is the time-complexity of computing the matrix-vector product Tx ?

18. (10 points) A real-valued univariate polynomial of degree n

$$p(x) = a_n x^n + a_{n-1} x^{n-1} + \cdots + a_1 x + a_0$$

can be represented using the following data structure

```
struct polynomial {  
    unsigned int n;  
    double *coef;  
};
```

where `coef` points to the first element of an array that stores the $n + 1$ coefficients (a_0, a_1, \dots, a_n) . You may assume that $a_n \neq 0$.

- (a) The value of the polynomial p at x can be evaluated using Horner's method: initialize $b_n := a_n$ and compute

$$\begin{aligned} b_{n-1} &:= a_{n-1} + b_n x \\ b_{n-2} &:= a_{n-2} + b_{n-1} x \\ &\vdots \\ b_0 &:= a_0 + b_1 x \end{aligned}$$

which yields $b_0 = p(x)$. Write a function that evaluates $p(x)$ using Horner's method.

Your function should have the following prototype:

```
double poly_eval(struct polynomial *p, double x);
```

- (b) The product

$$q(x) = p_1(x)p_2(x)$$

of two polynomials $p_1(x)$ and $p_2(x)$ of degree n_1 and n_2 , respectively, is itself a polynomial $q(x)$, and the degree of $q(x)$ is $n_1 + n_2$. Write a function that takes two polynomials p_1 and p_2 as inputs and computes q .

Your function should have the following prototype:

```
struct polynomial * poly_mul(  
    const struct polynomial *p1,  
    const struct polynomial *p2  
);
```