

Week 2 solutions

Exercises — Part I

Exercise 1

Both statements are false.

Exercise 2: Do exercises 3-1 and 4-1 in “Beginning C”

BC exercise 3-1

```
/*
    02635 Mathematical Software Programming
    Beginning C, exercise 3-1
*/
#include <stdio.h>

int main(void) {

    int choice=0;
    float temperature=0.0;

    printf("Temperature conversion"
           "\n- please select one of the following options:\n");
    printf("  1. Convert from degrees Celcius to degrees Fahrenheit\n");
    printf("  2. Convert from degrees Fahrenheit to degrees Celcius\n");
    printf("\nEnter your choice [1 or 2]: ");
    scanf("%i",&choice);

    if (choice == 1) {
        printf("Please enter temperature in degrees Celcius: ");
        scanf("%f",&temperature);
        printf("Temperature in degrees Fahrenheit: %.1f F\n",
               temperature*1.8+32);
    }
    else if (choice == 2) {
        printf("Please enter temperature in degrees Fahrenheit: ");
        scanf("%f",&temperature);
        printf("Temperature in degrees Celcius: %.1f C\n",
               (temperature-32)/1.8);
    }
    else {
        printf("Invalid choice.\n");
        return -1;
    }
}
```

```

    return 0;
}

```

BC exercise 4-1

```

/*
    02635 Mathematical Software Programming
    Beginning C, exercise 4-1
*/
#include <stdio.h>

int main(void) {

    int size=0;
    printf("Multiplication table - please enter size: ");
    scanf("%d",&size);

    // print first line with integers
    printf("      | ");
    for (int j=1;j<=size;j++)
        printf("%5d ",j);
    printf("\n-----");
    for (int j=1;j<=size;j++)
        printf("-----");
    printf("\n");

    // print table
    for (int i=1;i<=size;i++) {
        printf("%5d | ",i);
        for (int j=1;j<=size;j++) {
            printf("%5d ",i*j);
        }
        printf("\n");
    }

    return 0;
}

```

Example output

```

Multiplication table - please enter size: 8
      |      1      2      3      4      5      6      7      8
-----
    1 |      1      2      3      4      5      6      7      8
    2 |      2      4      6      8     10     12     14     16

```

3	3	6	9	12	15	18	21	24
4	4	8	12	16	20	24	28	32
5	5	10	15	20	25	30	35	40
6	6	12	18	24	30	36	42	48
7	7	14	21	28	35	42	49	56
8	8	16	24	32	40	48	56	64

Exercise 3: Do exercises 2, 3, and 4 (p. 39) in “Writing Scientific Software”

WSS exercise 2

```

/*
    02635 Mathematical Software Programming
    Writing Scientific Software, exercise 2
*/
#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include <fenv.h>
#ifndef RMODE
#define RMODE FE_TONEAREST
#endif

#pragma STDC FENV_ACCESS ON
int main(void) {
    double x = 1.0;

    if (fesetround(RMODE)) {
        fprintf(stderr, "Failed to set rounding mode\n");
        exit(EXIT_FAILURE);
    }
    for (int k=0;k<=16;k++) {
        printf("f(10^(%-3d)) = %.16le\n", -k, (1-cos(x))/(x*x));
        x *= 0.1;
    }
    return 0;
}

```

Output

```

f(10^(0 )) = 4.5969769413186023e-01
f(10^(-1 )) = 4.9958347219742893e-01
f(10^(-2 )) = 4.9999583334736619e-01
f(10^(-3 )) = 4.9999995832550309e-01
f(10^(-4 )) = 4.9999999696126418e-01
f(10^(-5 )) = 5.0000004137018506e-01

```

```

f(10^(-6)) = 5.0004445029117006e-01
f(10^(-7)) = 4.9960036108131994e-01
f(10^(-8)) = 0.0000000000000000e+00
f(10^(-9)) = 0.0000000000000000e+00
f(10^(-10)) = 0.0000000000000000e+00
f(10^(-11)) = 0.0000000000000000e+00
f(10^(-12)) = 0.0000000000000000e+00
f(10^(-13)) = 0.0000000000000000e+00
f(10^(-14)) = 0.0000000000000000e+00
f(10^(-15)) = 0.0000000000000000e+00
f(10^(-16)) = 0.0000000000000000e+00

```

The limit of $f(x) = \frac{1-\cos(x)}{x^2}$ as $x \rightarrow 0$ can be found using L'Hôpital's rule, i.e.,

$$\lim_{x \rightarrow 0} f(x) = \lim_{x \rightarrow 0} \frac{\sin(x)}{2x} = \lim_{x \rightarrow 0} \frac{\cos(x)}{2} = \frac{1}{2},$$

and hence the computed value of $f(x)$ for small x is quite far off. The problem is catastrophic cancellation in the numerator of $f(x)$.

The cosine function can be represented by the following Taylor series

$$\cos(x) = \sum_{k=0}^{\infty} (-1)^k \frac{x^{2k}}{(2k)!}.$$

If x is close to zero, the fourth-order approximation

$$\cos(x) \approx 1 - \frac{1}{2}x^2 + \frac{1}{24}x^4$$

is reasonably accurate. Indeed, for $|x| \leq 10^{-4}$ the magnitude of the fourth-order term is smaller than 0.02ϵ (assuming double precision floating-point arithmetic). With this approximation we have

$$f(x) = \frac{1 - \cos(x)}{x^2} \approx \frac{1}{2} - \frac{1}{24}x^2$$

for x sufficiently close to zero. For $x = 10^{-6}$ we have

$$\begin{aligned} \text{fl}(1 - \text{fl}(\cos(x))) &\approx \left(1 - \left(1 - \frac{1}{2}x^2 + \frac{1}{24}x^4\right)(1 + \delta_1)\right)(1 + \delta_2) \\ &\approx \frac{10^{-12}}{2} - \frac{10^{-24}}{24} - \delta_1 \end{aligned}$$

where $|\delta_1| \leq u$ and $|\delta_2| \leq u$, and hence the relative error is approximately

$$e_{\text{rel}} \approx \frac{|\frac{1}{24} \cdot 10^{-24} + \delta_1|}{|\frac{1}{2} \cdot 10^{-12}|} \lesssim \frac{2 \cdot 10^{-16}}{10^{-12}} = 2 \cdot 10^{-4}.$$

The computed value for $x = 10^{-6}$ yields a numerator with the (approximate) absolute error

$$e_{\text{abs}} \approx |0.500044 - 0.5 + \frac{1}{24} \cdot 10^{-12}| \cdot 10^{-12} \approx 4.4 \cdot 10^{-17}$$

which corresponds to approximately $0.4u$, and hence the relative error in the numerator is approximately

$$e_{\text{rel}} = \frac{e_{\text{abs}}}{|1 - \cos(x)|} \approx \frac{4.4 \cdot 10^{-17}}{\frac{1}{2} \cdot 10^{-12}} \approx 8.8 \cdot 10^{-5}.$$

WSS exercise 3

```

/*
    02635 Mathematical Software Programming
    Writing Scientific Software, exercise 3
*/
#include <stdio.h>
#include <math.h>

int main(void) {

    double x,y,xt,yt,tmp;

    printf("Input x: ");
    scanf("%lf",&x);
    printf("Input y: ");
    scanf("%lf",&y);

    xt = fabs(x);
    yt = fabs(y);

    // ensure that yt >= xt by swapping xt and yt if xt > yt
    if (xt > yt) {
        tmp = yt;
        yt = xt;
        xt = tmp;
    }

    if ( (yt != 0.0) && (yt != INFINITY) ) {
        tmp = xt/yt;
        printf("sqrt(x^2 + y^2) = %le\n", fabs(yt)*sqrt(1.0+tmp*tmp));
    }
    else {
        printf("sqrt(x^2 + y^2) = %le\n", yt);
    }

    return 0;
}

```

WSS exercise 4

The solutions of a quadratic equation $ax^2 + bx + c = 0$ are

$$x = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}.$$

If b is positive and b^2 is **large** compared to ac , then $-b + \sqrt{b^2 - 4ac}$ is likely to result in

catastrophic cancellation. If x_1 and x_2 denote the two solutions, it is easy to verify that

$$x_1 x_2 = -\frac{c}{a},$$

as stated in the exercise. Thus, if b^2 is large compared to $4ac$, we compute

$$x_1 = \frac{-b - \operatorname{sgn}(b)\sqrt{b^2 - 4ac}}{2a}, \quad x_2 = \frac{c}{ax_1}.$$

Remark: The determinant, $b^2 - 4ac$, may also be subject to cancellation, and hence it is generally a good idea to compute this in extended precision.

```

/*
    02635 Mathematical Software Programming
    Writing Scientific Software, exercise 4
*/
#include <stdio.h>
#include <math.h>

int main(void) {

    double a,b,c,xm,xp,det;

    printf("Solve quadratic equation a*x^2 + b*x + c == 0\n\n");

    // prompt user to enter a,b,c
    printf("Input a: ");
    scanf("%lf",&a);
    printf("Input b: ");
    scanf("%lf",&b);
    printf("Input c: ");
    scanf("%lf",&c);

    det = b*b - 4*a*c;

    if (a == 0) {
        if (b != 0)
            printf("x = %.4le\n",-c/b);
        else
            printf("a and b are both zero.\n");
        return 0;
    }
    else { // a is nonzero
        if (det < 0) { // complex roots
            printf("Complex roots\n");
            printf("x1 = %.4le + i*%.4le, ",-b/(2*a),sqrt(-det)/(2*a));
            printf("x2 = %.4le - i*%.4le\n",-b/(2*a),sqrt(-det)/(2*a));
            return 0;
        }
    }
}

```

```

else if (b*b > 10*a*c && b > 0) { // special case 1
    xm = -b - sqrt(det);
    xm /= 2*a;
    xp = c/(a*xm);
}
else if (b*b > 10*a*c && b < 0) { // special case 2
    xp = -b + sqrt(det);
    xp /= 2*a;
    xm = c/(a*xp);
}
else { // default case
    xp = (-b + sqrt(det))/(2*a);
    xm = (-b - sqrt(det))/(2*a);
}
printf("Real roots\n");
printf("x1 = %.4le, x2 = %.4le\n",xp,xm);
return 0;
}
}

```

Exercise 4: Modify your code from exercise 2 in “Writing Scientific Software” so that it uses the “round toward zero” rounding mode.

Changing the rounding mode does not mitigate catastrophic cancellation. However, it reveals that some results (e.g., at $x = 10^{-7}$) are very sensitive to rounding errors:

```
f(10^(-7 )) = 5.1070259132757200e-01
```

Exercise 5: Table with floating-point operations

Write a program that prints a table with values $x \text{ op } y$ where

- x and y are `-INFINITY`, `-1.0`, `-0.0`, `0.0`, `1.0`, `INFINITY`, or `NAN`
- op is one of the arithmetic operators `*`, `/`, `+`, or `-`, or one of the relational operators `==`, `!=`, `>`, or `<`

```

#include <stdio.h>
#include <math.h>

#define xstr(s) str(s)
#define str(s) #s
#ifdef op
#define op ==
#endif

```

```

int main(void) {
    int i,j;
    double arr[7] = {-INFINITY,-1.0,-0.0,0.0,1.0,INFINITY,NAN};

    printf("op: %2s | ", xstr(op));
    for (i=0;i<7;i++)
        printf("%6.1f ", arr[i]);
    putchar('\n');
    for (i=0;i<57;i++)
        putchar('-');
    putchar('\n');

    for (i=0;i<7;i++) {
        printf("%6.1f | ", arr[i]);
        for (j=0;j<7;j++) {
            printf("%6.0f ", (double) (arr[i] op arr[j]));
        }
        printf("\n");
    }
    return 0;
}

```

```
$ gcc -Wall -lm -D op='*' fpops.c -o fpops
```

```
$ ./fpops
```

```

op:  * |   -inf   -1.0   -0.0    0.0    1.0    inf    nan
-----
- inf |    inf    inf    nan    nan   -inf   -inf    nan
-1.0 |    inf     1     0    -0    -1   -inf    nan
-0.0 |    nan     0     0    -0    -0    nan    nan
 0.0 |    nan    -0    -0     0     0    nan    nan
 1.0 |   -inf    -1    -0     0     1    inf    nan
 inf |   -inf   -inf    nan    nan    inf    inf    nan
 nan |    nan    nan    nan    nan    nan    nan    nan

```

Exercises — Part II

Exercise 8

```

#include <stdio.h>
#include <math.h>

#define RECTANGLE 1
#define TRAPEZOIDAL 2

int main(void) {

```



```
double a, b, h, val = 0, x;
int n, method;

// Print welcome message
printf("Computes an approximation of the definite integral\n\n"
      "   $\int_a^b \exp(-x^2) dx$ \n\n"
      "using numerical integration.\n\n");

// Prompt user to enter integration limits
printf("Please enter the integration limit a: ");
scanf("%lf", &a);
printf("Please enter the integration limit b: ");
scanf("%lf", &b);
// Check that a < b
if (a >= b) {
    printf("error: a must be less than b\n");
    return -1;
}

// Prompt user to enter number of subintervals
printf("Please enter the number of subintervals: ");
scanf("%d", &n);
// Check that n is positive
if (n <= 0) {
    printf("error: n must be positive\n");
    return -1;
}

// Prompt user to choose method
printf("Please select integration rule (%i. rectangle rule,"
      " %i. trapezoidal rule): ", RECTANGLE, TRAPEZOIDAL);
scanf("%d", &method);
// Check user input
if (!(method == RECTANGLE || (method == TRAPEZOIDAL))) {
    printf("error: unknown method\n");
    return -1;
}

// Compute approximation to definite integral and print result
h = (b-a)/n;
if (method == RECTANGLE) {
    for (int i=0; i<n; i++) {
        x = a + (i+0.5)*h;
        val += h*exp(-x*x);
    }
}
else if (method == TRAPEZOIDAL) {
```

```
    val = 0.5*h*(exp(-a*a) + exp(-b*b));  
    for (int i=1; i<n-1; i++) {  
        x = a+i*h;  
        val += h*exp(-x*x);  
    }  
}  
printf("Approximate value of definite integral: %.8e\n", val);  
  
return 0;  
}
```

Exercise 9

```
#include <stdio.h>  
#include <math.h>  
  
int main(void) {  
  
    double a, b, h, val1, val2, x;  
    int n;  
  
    // Print welcome message  
    printf("Computes an approximation of the definite integral\n\n"  
           "  int_a^b exp(-x^2) dx\n\n"  
           "using numerical integration.\n\n");  
  
    // Prompt user to enter integration limits  
    printf("Please enter the integration limit a: ");  
    scanf("%lf", &a);  
    printf("Please enter the integration limit b: ");  
    scanf("%lf", &b);  
    // Check that a < b  
    if (a>=b) {  
        printf("error: a must be less than b\n");  
        return -1;  
    }  
  
    // Prompt user to enter number of subintervals  
    printf("Please enter the number of subintervals: ");  
    scanf("%d",&n);  
    // Check that n is positive  
    if (n <= 0) {  
        printf("error: n must be positive\n");  
        return -1;  
    }  
  
    // Compute results and print table
```

```

printf("Parameters:\n\n a = %.3e\n b = %.3e\n n = %i\n\n",a,b,n);
printf("Results:\n\n");
printf("%3s %-14s %-14s\n","n","Rectangle","Trapezoidal");
printf("-----\n");
for (int i=1;i<=n;i++) {
    h = (b-a)/i;

    // Rectangle rule
    val1 = 0.0;
    for (int j=0; j<i; j++) {
        x = a + (j+0.5)*h;
        val1 += h*exp(-x*x);
    }

    // Trapezoidal rule
    val2 = 0.5*h*(exp(-a*a) + exp(-b*b));
    for (int j=1; j<i; j++) {
        x = a+j*h;
        val2 += h*exp(-x*x);
    }

    // Print row
    printf("%3i %.8e %.8e\n",i,val1,val2);
}

return 0;
}

```

Optional exercise: Monto Carlo integration

```

#include <stdlib.h>
#include <stdio.h>
#include <math.h>
#include <time.h>

int main(void) {

    double a, b, val, u;
    int n;

    // Initialize random number generator
    srand(time(NULL));

    // Print welcome message
    printf("Computes an approximation of the definite integral\n\n"
           "  int_a^b exp(-x^2) dx\n\n"
           "using Monte Carlo integration.\n\n");
}

```

```
// Prompt user to enter integration limits
printf("Please enter the integration limit a: ");
scanf("%lf", &a);
printf("Please enter the integration limit b: ");
scanf("%lf", &b);
// Check that a < b
if (a>=b) {
    printf("error: a must be less than b\n");
    return -1;
}

// Prompt user to enter number of samples
printf("Please enter the number of samples: ");
scanf("%d",&n);
// Check that n is positive
if (n <= 0) {
    printf("error: n must be positive\n");
    return -1;
}

// Compute result
val = 0.0;
for (int i=1;i<=n;i++) {
    u = a + (b-a)*rand()/RAND_MAX;
    val = (1.0-1.0/i)*val + exp(-u*u)/i;
}
printf("Approximate value of definite integral: %.8e\n", val*(b-a));

return 0;
}
```