

Week 7 solutions

Exercises

Part I: Timing datasize1()

```

/* ex1.c */
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define MAX_SIZE 16777216 // 128*1024*1024/8 elements
extern int datasize1(int);
double arr[MAX_SIZE]; // global array of length MAX_SIZE

#define mytimer clock
#define delta_t(a,b) (1.0e3 * ((b) - (a)) / CLOCKS_PER_SEC)
#define MIN_RUNTIME 2000 // run iterations for at least MIN_RUNTIME msecs

int main(int argc, char *argv[]) {
    clock_t t1, t2;
    double tcpu;
    int mem_acc;
    int iter;

    printf("# Testing function datasize1:\n");
    printf("%10s %9s\n", "Mem (kB)", "Mflop/s");
    for(size_t i = 2048; i <= MAX_SIZE; i *= 2) {
        tcpu = 0.0; iter = 0;
        t1 = mytimer();
        do {
            mem_acc = datasize1(i);
            t2 = mytimer();
            tcpu = delta_t(t1, t2);
            iter++;
        } while (tcpu < MIN_RUNTIME);

        // Print memory (kB) and Mflop/s
        printf("%10.0f %9.2e\n", 8.0*i/1024, (1e-3*iter)*mem_acc/tcpu);
    }

    return(0);
}

```

```

/* datasize1.c */
extern double arr[];

```

```
int datasize1(int elem) {  
  
    for (int i=0; i<elem; i++)  
        arr[i] *= 3;  
  
    return(elem);  
}
```

Part II

1. Implement the functions `my_dgemv_v1()` and `my_dgemv_v2()`:

```
void my_dgemv_v1(  
    int m,           /* number of rows */  
    int n,           /* number of columns */  
    double alpha,    /* scalar */  
    double ** A,     /* two-dim. array A of size m-by-n */  
    double * x,      /* one-dim. array x of length n */  
    double beta,     /* scalar */  
    double * y       /* one-dim. array x of length m */  
) {  
    int i,j;  
    for (i=0;i<m;i++) {  
        y[i] *= beta;  
        for (j=0;j<n;j++) {  
            y[i] += alpha*A[i][j]*x[j];  
        }  
    }  
    return;  
}
```

```
void my_dgemv_v2(  
    int m,           /* number of rows */  
    int n,           /* number of columns */  
    double alpha,    /* scalar */  
    double ** A,     /* two-dim. array A of size m-by-n */  
    double * x,      /* one-dim. array x of length n */  
    double beta,     /* scalar */  
    double * y       /* one-dim. array x of length m */  
) {  
    int i,j;  
    for (i=0;i<m;i++)  
        y[i] *= beta;  
    for (j=0;j<n;j++) {  
        for (i=0;i<m;i++) {  
            y[i] += alpha*A[i][j]*x[j];  
        }  
    }  
}
```

```

    }
}
return;
}

```

2. See solution to exercise 4 below.
3. See solution to exercise 4 below.
4. Repeat the two timing experiments with compiler optimizations:

```

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include "my_dgemv_v1.h"
#include "my_dgemv_v2.h"
#include "array_2d.h"

#define NREPEAT 200
#define mytimer clock
#define delta_t(a,b) (1.0e3 * ((b) - (a)) / CLOCKS_PER_SEC)

int main(int argc, char *argv[]) {

    int i, m, n, N = NREPEAT;
    double *x, *y, **A, tcu1, tcu2;
    clock_t t1, t2;

    printf("%4s %8s %8s\n", "n", "v1", "v2");
    for (m = 100; m <= 1000; m += 100) {
        n = m;

        /* Allocate memory */
        A = malloc_2d(m, n);
        x = malloc(n*sizeof(*x));
        y = malloc(m*sizeof(*y));
        if (A == NULL || x == NULL || y == NULL) {
            fprintf(stderr, "Memory allocation error..\n");
            free_2d(A); free(x); free(y);
            exit(EXIT_FAILURE);
        }

        /* CPU time for my_dgemv_v1 */
        t1 = mytimer();
        for (i = 0; i < N; i++)
            my_dgemv_v1(m, n, 1.0, A, x, 0.0, y);
        t2 = mytimer();
        tcu1 = delta_t(t1,t2)/N;
    }
}

```

```

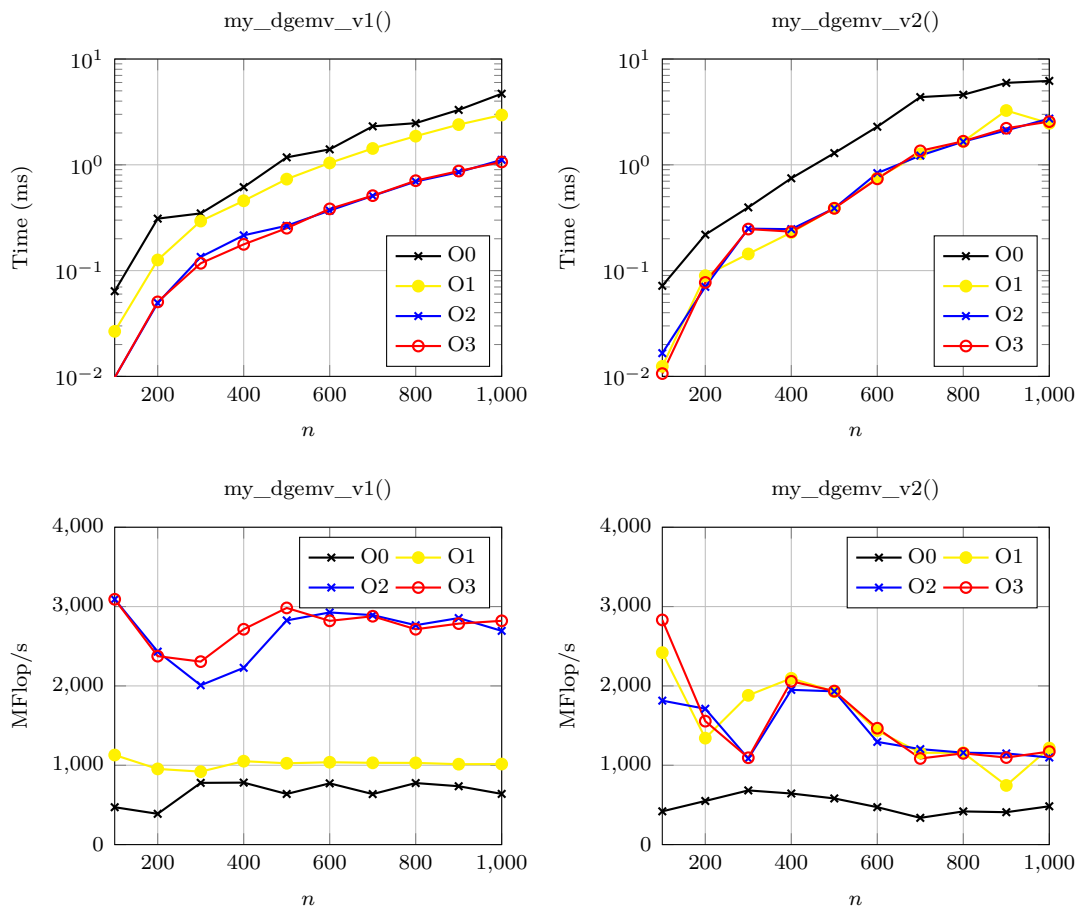
/* CPU time for my_dgemv_v2 */
t1 = mytimer();
for (i = 0; i < N; i++)
    my_dgemv_v2(m, n, 1.0, A, x, 0.0, y);
t2 = mytimer();
tcpu2 = delta_t(t1,t2)/N;

/* Print n and results */
printf("%4d %8.3e %8.3e\n", n, tcpu1, tcpu2);

/* Free memory */
free_2d(A);
free(x);
free(y);
}

return EXIT_SUCCESS;
}

```



The results show that both versions benefit quite a bit from compiler optimization (for all n). Moreover, `my_dgemv_v1` is significantly faster than `my_dgemv_v2` when n is large. Indeed, the spacial locality is much better in the first variant of the method since it accesses the elements of A row-by-row in accordance with the row-major storage.

Note that the CPU times may differ (significantly) on other systems.

5. The above implementations of `my_dgemv_v1` and `my_dgemv_v2` both require $3mn+m$ floating-point operations for a matrix A of size $m \times n$. Assuming that the time is T seconds and $m = n$, we can compute the performance in MFlop/s as

$$\frac{3n^2 + n}{10^6 \cdot T}.$$

Finally, note that nested loops in `my_dgemv_v1` can also be implemented as follows:

```
for (i=0;i<m;i++) {  
    double dotx = 0.0;  
    for (j=0;j<n;j++) {  
        dotx += A[i][j]*x[j];  
    }  
    y[i] = alpha*dotx + beta*y[i];  
}
```

This implementation requires $2mn + 3m$ Flops instead of $3mn + m$ Flops. Similarly, `my_dgemv_v2` can also be implemented with a lower Flop count than that of the implementation provided above.